Donghu Kim

quagmire@kaist.ac.kr | i-am-proto.github.io | I-AM-PROTO

Research Interest

Efficient Reinforcement Learning

- Making RL work with as little samples and/or compute as possible.
- Related: Simba, SimbaV2, AtariPB

Plasticity

- Maintaining plasticity (the ability to train) when the data distribution is constantly shifting/expanding.
- Related: Dynamic MoE, Catastrophic Interference, Hare&Tortoise

Education

KAIST — *Seongnam, Korea*
M.S. Student in AI (GPA: 3.8/4.3, Advisor: Jaegul Choo) — *Mar. 2024 - Present*

Korea University — *Seoul, Korea*
B.S. in Computer Science (Major GPA: 4.5/4.5, Cumulative GPA: 4.2/4.5) — *Mar. 2018 - Feb. 2024*

Work Experience

Krafton AI — *Seoul, Korea*
Research Intern — *June. 2025 - Present*

- Physical Intelligence Team (TBD).

Publications & Preprints

SimbaV2: Hyperspherical Normalization for Scalable Deep Reinforcement Learning — *ICML'25*
Hojoon Lee*, Youngdo Lee*, Takuma Seno, **Donghu Kim**, Peter Stone, Jaegul Choo — ***Spotlight***

- arXiv / project page / code

SimBa: Simplicity Bias for Scaling Up Parameters in Deep Reinforcement Learning — *ICLR'25*
Hojoon Lee*, Dongyoon Hwang*, **Donghu Kim**, ... , Jaegul Choo, Peter Stone, Takuma Seno — ***Spotlight***

- arXiv / project page / code

Do's and Don'ts: Learning Desirable Skills with Instruction Videos — *NeurIPS'24*
Hyunseung Kim, Byungkun Lee, Hojoon Lee, Dongyoon Hwang, **Donghu Kim**, Jaegul Choo — *Poster*

- arXiv / project page

ATARI-PB: Investigating Pre-Training Objectives for Generalization in Pixel-Based RL — *ICML'24*
Donghu Kim*, Hojoon Lee*, Kyungmin Lee*, Dongyoon Hwang, Jaegul Choo — *Poster*

- arXiv / project page / code

Slow and Steady Wins the Race: Maintaining Plasticity with Hare and Tortoise Networks — *ICML'24*
Hojoon Lee, Hyeonseo Cho, Hyunseung Kim, **Donghu Kim**, Dugki Min, Jaegul Choo, Clare Lyle — *Poster*

- arXiv / code

Projects

Dynamic Mixture-of-Experts — *2025*

- Explored dynamically adding new experts to Mixture-of-Experts layers to maintain plasticity in environments with severe distribution shifts (e.g., Craftax).
- report / slides

KAN RL *2024*

- Implemented Kolmogorov-Arnold Network in sequential Atari environments and investigated its relevance to catastrophic forgetting and plasticity.
- report (Colab)

RL Basic Tutorial *2024*

- Developed and delivered a series of three lectures on reinforcement learning for a government-funded bootcamp program in Korea.
- page / material1 (Korean) / material2 (Korean)

Character-level BERT *2022*

- Proposed a character-level tokenizer for BERT to enhance robustness against character-level attacks common in spam emails.
- report / code

Honors & Awards

Korea University	Academic Excellence Award	*2019, 2022*
NCsoft AI Fellowship	Starcraft AI Competition Silver Prize ($2000)	*2019*
Korea Student Aid Foundation	Presidential Science Scholarship (Total $40000)	*2018*